



构建高性能MySQL之索引

运维与系统开发部DBA 王鹏

目录

Contents

第一章：慢SQL引发的思考

第二章：索引基础理论与数据结构

第三章：理解InnoDB索引与MyISAM索引



慢SQL引发的思考



```

SELECT presentid, count(*) AS cnt, sum(amount) AS amount
  FROM presentorders
 WHERE 1 = 1
    AND ((presentid IN ('772378151', '772378704')) OR amount > 0)
    AND STATUS = '1'
    AND ctime >= '2015-08-28'
    AND ctime < '2015-08-29'
 GROUP BY presentid
HAVING cnt >= '2'
 ORDER BY cnt DESC;

```

```

PRIMARY KEY (`id`),
KEY `idx_presentid_ctime` (`presentid`, `ctime`),
KEY `idx_patientid_ctime` (`patientid`, `ctime`),
KEY `idx_spaceid_ctime` (`spaceid`, `ctime`),
KEY `idx_spaceid_patientid` (`spaceid`, `patientid`),
KEY `idx_ptime_space` (`ptime`, `spaceid`),
KEY `idx_userid` (`userid`),
KEY `idx_contractid` (`contractid`),
KEY `idx_space_present` (`spaceid`, `presentid`)

```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	presentorders	index	idx_presentid_ctime,idx_space_present	idx_presentid_ctime	8	NULL	1293646	Using where; Using temporary; Using filesort


```
mysql> SELECT presentid, count(*) AS cnt, sum(amount) AS amount FROM presentorders WHERE 1 = 1 AND ((presentid IN ('772378151', '772378704')) OR amount > 0) AND STATUS = '1' AND ctime >= '2015-08-28' AND ctime < '2015-08-29' GROUP BY presentid HAVING cnt >= '2' ORDER BY cnt DESC;
```

presentid	cnt	amount
3210338199	266	266.00
3210336341	111	111.00
3210353495	100	1000.00
3210350260	66	660.00
3216496931	32	32.00
3210376439	28	840.00
3210381066	27	810.00
3128226885	26	2600.00
3097882366	15	1500.00
3216506320	14	140.00
3210329315	12	12.00
3210348087	9	90.00
2254312142	6	6.00
3216509041	5	150.00
3132705573	3	3.00
699322060	3	60.00
699322045	2	40.00
800871576	2	60.00
1004130974	2	20.00
2254329894	2	60.00
3210371155	2	60.00
3097644213	2	200.00

22 rows in set (3.61 sec)

```
mysql> select count(*) from presentorders where presentid in ('772378151','772378704');
```

count(*)

349

1 row in set (0.00 sec)

```
mysql> select count(*) from presentorders where ctime >= '2015-08-28'  
-> AND ctime < '2015-08-29';
```

count(*)

2115

1 row in set (0.01 sec)

```
mysql> SELECT presentid, count(*) AS cnt, sum(amount) AS amount FROM presentorders WHERE 1 = 1 AND ((presentid IN ('772378151', '772378704')) OR amount > 0) AND STATUS = '1' AND ctime >= '2015-08-28' AND ctime < '2015-8-29' GROUP BY presentid HAVING cnt >= '2' ORDER BY cnt DESC;
```

presentid	cnt	amount
3210338199	266	266.00
3210336341	111	111.00
3210353495	100	1000.00
3210350260	66	660.00
3216496931	32	32.00
3210376439	28	840.00
3210381066	27	810.00
3128226885	26	2600.00
3097882366	15	1500.00
3216506320	14	140.00
3210329315	12	12.00
3210348087	9	90.00
2254312142	6	6.00
3216509041	5	150.00
699322060	3	60.00
3132705573	3	3.00
800871576	2	60.00
2254329894	2	60.00
3210371155	2	60.00
1004130974	2	20.00
699322045	2	40.00
3097644213	2	200.00

22 rows in set (0.01 sec)

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	presentorders	range	idx_presentid_ctime,idx_space_present,idx_c	idx_c	4	NULL	2113	Using index condition; Using where; Using temporary; Using filesort

1 row in set (0.00 sec)

Table	Non_unique	Key_name	Seq_in_index	column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
presentorders	0	PRIMARY	1	id	A	1294095	NULL	NULL		BTREE
presentorders	1	idx_presentid_ctime	1	presentid	A	510	NULL	NULL		BTREE
presentorders	1	idx_presentid_ctime	2	ctime	A	1294095	NULL	NULL		BTREE


```
mysql> select distinct a.userid as userid
->      from doctoroutlines a
->      inner join tab_DiseaseDoctor b on b.fld_DoctorId = a.doctorid and b.fld_DiseaseDoctorGoodVoteCount >0
->      where b.fld_DiseaseId = '27'
->      and a.userid not in ('93546','42395','945963')order by b.fld_DiseaseDoctorRank desc limit 1;
```

```
+-----+
| userid |
+-----+
| 421234 |
+-----+
1 row in set (1.03 sec)
```

table	type	key	key_len	ref	rows	Extra
b	index	fld_DiseaseDoctorRank	4	NULL	218	Using where; Using temporary
a	ref	doctorid_index	4	db_Hdf.b.fld_DoctorId	1	Using where

```
doctoroutlines  KEY `doctorid_index` (`doctorid`)
tab_DiseaseDoctor KEY `fld_DiseaseDoctorRank` (`fld_DiseaseDoctorRank`)
tab_DiseaseDoctor KEY `fld_DiseaseId` (`fld_DiseaseId`)
```



```
mysql> create index idx_c on tab_DiseaseDoctor(fld_DiseaseId,fld_DiseaseDoctorRank);
mysql> select distinct a.userid as userid
->      from doctoroutlines a
->      inner join tab_DiseaseDoctor b on b.fld_DoctorId = a.doctorid and b.fld_DiseaseDoctorGoodVoteCount >0
->      where b.fld_DiseaseId = '27'
->      and a.userid not in ('93546','42395','945963')order by b.fld_DiseaseDoctorRank desc limit 1;
```

```
+-----+
| userid |
+-----+
| 421234 |
+-----+
```

1 row in set (0.02 sec)

table	type	key	key_len	ref	rows	Extra
b	ref	idx_c	4	const	2523	Using where; Using temporary
a	ref	doctorid_index	4	db_Hdf.b.fld_DoctorId	1	Using where

```
mysql> select distinct a.userid as userid
->   from doctoroutlines a,
->   (select b.fld_DoctorId
->       from tab_DiseaseDoctor b
->       where b.fld_DiseaseId = '27'
->       order by b.fld_DiseaseDoctorRank desc) c
->   where a.doctorid = c.fld_DoctorId
->   and a.userid not in ('93546', '42395', '945963') limit 1;
```

```
+-----+
| userid |
+-----+
| 421234 |
+-----+
```

1 row in set (0.01 sec)

table	type	key	key_len	ref	rows	Extra
<derived2>	ALL	NULL	NULL	NULL	2523	Using temporary
a	ref	doctorid_index	4	c.fld_DoctorId	1	Using where
b	ref	fld_DiseaseId	4	const	2523	Using where; Using filesort

```

select d.id as id
  from doctorpatientrefs d
 inner join flowcontentmarks e
    on d.id = e.doctorpatientrefid
 where d.userid <> '1293826'
    and d.status not in ('1', '3')
    and e.status = 'AUDIT'
    and e.mark in ('LEVEL_2', 'LEVEL_3')
    and d.spaceid in ('2498157',
                      '697669779',
                      '546515',
                      '1967854',
                      '104831',
                      '295212',
                      '41600074',
                      '2941372',
                      '2921923',
                      '1887133856');

```

```

| 2390058401 |
| 2390433464 |
| 2534292811 |
| 2868079996 |
+-----+

```

```

61 rows in set
~

```

(1.45 sec)

table	type	key	key_len	ref	rows	Extra
d	range	ix_spaceid_hosttype_status	4	NULL	580	Using index condition; Using where
e	eq_ref	ix_marks_ref	4	flow_avatar.d.id	1	Using where

```
mysql> select count(*) from doctorpatientrefs d where d.userid <> '1293826'
->    and d.status not in ('1', '3') and d.spaceid in ('2498157',
->                                                    '697669779',
->                                                    '546515',
->                                                    '1967854',
->                                                    '104831',
->                                                    '295212',
->                                                    '41600074',
->                                                    '2941372',
->                                                    '2921923',
->                                                    '1887133856');
```

```
+-----+
| count(*) |
+-----+
|    121661 |
+-----+
```

```
mysql> select count(*) from flowcontentmarks e where e.status = 'AUDIT'
->    and e.mark in ('LEVEL_2', 'LEVEL_3');
```

```
+-----+
| count(*) |
+-----+
|    27236 |
+-----+
```

```
mysql> select d.id as id
->          from flowcontentmarks e STRAIGHT_JOIN doctorpatientrefs d
->          on d.id = e.doctorpatientrefid
-> where d.userid <> '1293826'
->          and d.status not in ('1', '3')
->          and e.status = 'AUDIT'
->          and e.mark in ('LEVEL_2', 'LEVEL_3')
->          and d.spaceid in ('2498157',
->                             '697669779',
->                             '546515',
->                             '1967854',
->                             '104831',
->                             '295212',
->                             '41600074',
->                             '2941372',
->                             '2921923',
->                             '1887133856');
```

2368780644
2385524207
2390058401
2390433464
2534292811
2868079996

```
+-----+
61 rows in set (0.21 sec)
```



第二章：索引基础理论与数据结构



XĪNHUÁ ZÌDIǎN

新华字典

第 11 版



商務印書館

根节点

心部 (中) 母 (母) 示石 69

分支节点 1

心部

母 (母) 部

示部

示 (中) 部


石部

分支节点 2

五至八画

叶子节点

叶子节点内到数据块号的指针



seek time 5ms
latency time $1/120/2 = 4.17\text{ms}$
transmission time 0.1

binary search

1	2	3	4	5	6	7	8	9	10	11
6	12	15	18	22	25	28	35	46	58	60


low = 1 mid = 6 high = 11

1	2	3	4	5	6
6	12	15	18	22	25


↑ ↑ ↑
low=1 mid=3 high=5

1 2 3
6 12 15
↑ ↑
low=1 high=2
↑
mid=1

B-Tree



B+Tree




Select ...
From Customer A, Orders B
Where A.Cust_Id = B.Cust_Id
And A.Cust_Id = 3975;

使用一个 HASH 集群
散列函数

HASH键	集群键	数据
I/O		


使用 B 树索引



理解InnoDB索引与MyISAM索引

MyISAM


Row number	col1	col2
0	99	8
1	12	56
2	3000	62
...		
9997	18	8
9998	4700	13
9999	3	93



INNODB


Row number	col1	col2
0	99	8
1	12	56
2	3000	62
...		
9997	18	8
9998	4700	13
9999	3	93



- Primary key columns (col1)
- TID Transaction ID
- RP Rollback pointer
- Non-PK columns (col2)




INNODB

Row number	col1	col2
0	99	8
1	12	56
2	3000	62
...		
9997	18	8
9998	4700	13
9999	3	93






InnoDB (clustered) table layout




MyISAM (nonclustered) table layout


Sequential insertion into the page: each new record is inserted after the previous one




When the page is full, insertion continues in a new page



Inserting UUIDs: new records may be inserted between previously inserted records, forcing them to be moved



Pages that were filled and flushed to disk may have to be read again



*Only the first 13 characters of the UUID are shown

Table	Rows	Time (sec)	Index size (MB)
userinfo	1,000,000	137	342
userinfo_uuid	1,000,000	180	544
userinfo	3,000,000	1233	1036
userinfo_uuid	3,000,000	4525	1707



End



Thank
you

系统是由每一个细节组成

